Accurate on Average, Blind at the Fork

Anonymous Author(s)
Affiliation
Address
email

Abstract: Before a robot executes an action, it should know whether its own uncertainty could change the outcome. People make this judgment with an intuitive physics engine [1]: they run a few noisy simulations over an uncertain scene and call it precarious when the outcomes disagree. We build this check for robots as a goal-free *fork probe*: sample the robot's measured perception and execution error, roll the intended action chunk forward 64 times, and flag it when the futures split. With a physics engine and a real camera-based perception model, the probe catches 88% of forks on unseen layouts with 7% false alarms, reading outcomes from images alone. Without the engine the probe needs a world model, and there it fails: a token-space world model that predicts 3–4$\times$ better than a no-motion baseline catches only 18%, because its sampled uncertainty has the right size but not where futures branch; five further generative variants catch at most 3%, while sampling perception uncertainty into the start and simulating deterministically reaches 69%, a ceiling that more data, frame history and a DINO-WM transformer do not lift. Safe deployment needs world models with *branching fidelity*, which average-error benchmarks do not measure.

Keywords: World models, intuitive physics, uncertainty, safe deployment

1 Introduction

Manipulation policies, from diffusion policies to vision-language-action models, act by committing to short action chunks [2, 3, 4]. Deploying them safely requires a check before each chunk runs. Existing monitors address *epistemic* uncertainty, flagging inputs the model has not seen [5, 6]. We address a second kind that remains even for a perfect model: *physical* uncertainty. When a robot picks a block standing between two others, the same command can lift it cleanly or tip a neighbour depending on millimetres of execution error and perception error. Such a chunk is risky even if its nominal outcome is fine: whether it succeeds depends on errors the robot cannot control.

People handle this with an *intuitive physics engine* (IPE): a few approximate simulations over an uncertain scene, whose disagreement signals precariousness [1, 7, 8]. Our goal is to give robots the same ability, which needs two things: an honest model of the robot's own uncertainty, and a model to simulate with. Outside simulation the second must be learned, which gives the physical fidelity of world models a concrete, safety-relevant test: a world model is useful here only if its rollouts branch where the real dynamics branch, a property prediction error does not measure.

Contributions. (1) A *fork probe*: a pre-execution check that propagates the robot's measured perception and execution uncertainty through a rollout engine and flags splits, with no task labels, failure data or fitted thresholds (Sec. 2). (2) Validation with a physics engine and real camera perception on three unseen test sets, which sets the reference a world model must reach (Sec. 3.1). (3) *Branching fidelity* as a world-model evaluation, and two learned token-space world models that predict well on average yet fail it, with a diagnosis (Sec. 3.2).

Submitted to the 10th Conference on Robot Learning (CoRL 2026). Do not distribute.

2 Setup and the fork probe

Task. Three Jenga-like blocks stand upright in a row; a simulated Franka Panda (MuJoCo) must pick the middle (red) block without toppling its neighbours. At each chunk start a policy commits to an 8-step action chunk; ours come from recorded episodes, but the probe treats the policy as a black box and sees only the chunk, one camera image and proprioception.

Fragility. A chunk is *fragile* if the robot's uncertainty, its execution error and its error in perceiving the scene, can send the outcome into qualitatively different regimes. This needs no goal or failure label, which a monitor around a generalist policy cannot assume, and it isolates exactly the cases where intent no longer determines the outcome. For our task we measure it physically: a state is a *fork* if some but not all of 64 futures sampled from this uncertainty (2–62) topple a neighbouring block. A chunk that fails every time is not a fork and needs a task-level check.

Uncertainty model. Execution error comes from the robot's measured tracking residuals (median chunk maximum 3.2 mm). Perception is a frozen DINOv2 encoder [9] on one close-camera image with a keypoint head for the three blocks in play (2.5 mm held-out error); its error model is measured once, offline, as a zero-mean Gaussian fitted to its errors on held-out images with known poses (a full 9×9 position covariance capturing correlated errors, plus rotation covariances).

Fork probe and rollout engine (Fig. 1a). The probe imagines 64 versions of the chunk. Each starts from the perceived scene shifted by one sampled perception error and runs the chunk with one sampled execution error; a 30-step imagined hold then lets the outcome settle. It flags the chunk if, at hold steps 10 and 30, the endings (object positions or image tokens) projected on their first principal axis prefer two Gaussians to one by BIC, are well separated (Ashman's $D > 2$ [10]), have at least 2 members in the smaller group, and keep the same grouping; these are fixed statistical conventions, not fitted to labels. A variant draws 5 scene hypotheses and flags if at least 3 split. The probe works with any rollout engine. Our reference is MuJoCo with the three blocks placed at the sampled poses and everything else (block geometry, masses, friction) taken from the simulator, which a real robot does not have; the learned alternative is a world model (Sec. 3.2).

Ground truth. One real outcome cannot grade a monitor, since a fragile chunk may succeed by luck. An oracle therefore samples the 64 futures from the *true* scene in exact physics and labels the state a fork as above. Two independent oracle runs give the same label on 99% of states. Every evaluation uses unseen block layouts (about 60 screened forks plus 100 safe states) under a protocol frozen before the data existed.

3 Results

3.1 Propagating uncertainty through a physics engine works

On the 112 states used to design the split test (exact physics, true scene), it flags all 23 topple forks and 10% of the 89 states without one. With real perception, perception uncertainty must be propagated: errors of 2–4 mm are comparable to execution noise, and a probe that trusts the estimate catches only 65–67% of forks, while the fork probe catches 77–91% with 7–17% false alarms (Table 1). Reading outcomes from DINOv2 tokens of rendered endings instead of object positions costs almost nothing, so the probe works from images alone: on test set C it catches 88% of forks (95% CI 77–94%) with 7% false alarms. This is the reference a world model must reach.

3.2 Propagating uncertainty through learned world models fails

Without privileged geometry a learned model must take MuJoCo's role, and the probe becomes its test: given one image and the 64 perturbed chunks, do its rollouts branch where physics branches? Our model compresses DINOv2 tokens by a fixed per-patch PCA to a $28 \times 28 \times 8$ latent (on held-out

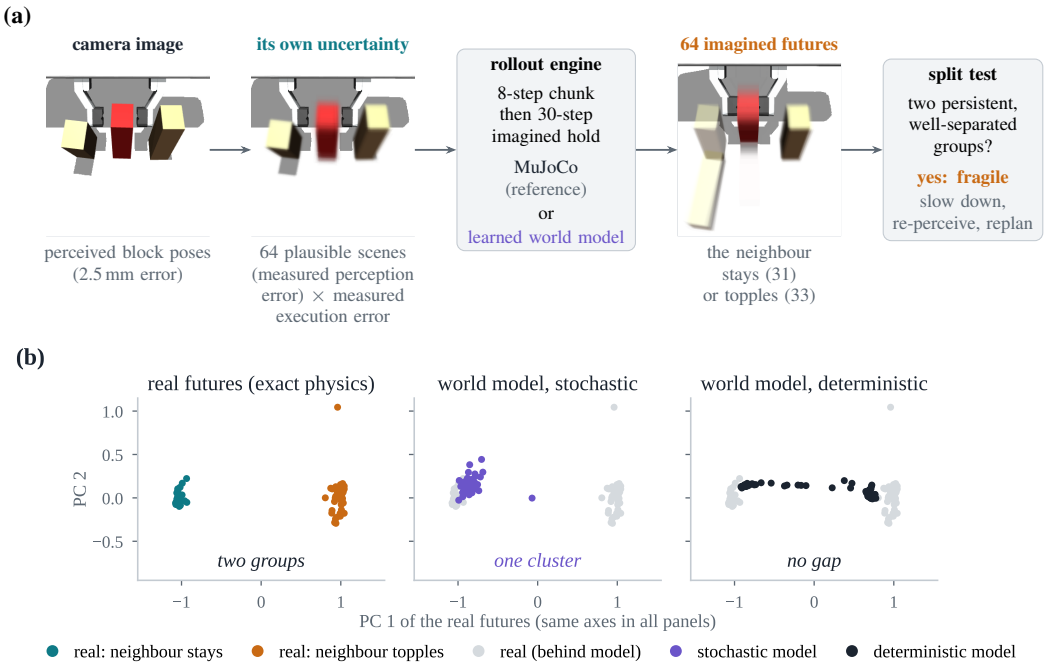


Figure 1: (a) The fork probe on one held-out state (exact renders; images average all 64 samples). (b) The same 64 futures in the world model's latent space, on the real futures' principal axes.

Table 1: Fork probe (MuJoCo, real perception) on unseen test sets: forks caught / false alarms (%).

outcomes read from object positions	Set A (153 states) caught	Set A (153 states) false	Set B (160) caught	Set B (160) false	Set C (156) caught	Set C (156) false
perception uncertainty ignored	65	23	65	12	67	24
fork probe	90	12	77	8	91	10
fork probe, 5 hypotheses	90	17	90	13	89	11
outcomes read from images: fork probe	88	12	77	8	88	9
5 hypotheses	–	–	–	–	**88**	**7**

real futures the split test on this latent catches 97% of forks) and predicts the next latent from the current one and the action with a flow-matching transition, drawing fresh noise at every step. It is trained on 8 futures per image whose start scenes were sampled within the measured perception error, so its target spread is the robot's own uncertainty; a deterministic version of the same network is the control. Both beat a no-motion baseline on 19-step latent error by 3–4×, yet on test set C the stochastic model catches 18% of forks and the deterministic one 60%, against 88% for the engine (Table 2).

Where the uncertainty went wrong. Compared with real futures (Table 2: total spread, share on the main axis, ratio on forks vs other states), the stochastic model has the right *amount* of uncertainty but not its *shape* or *placement*: the spread is isotropic and nearly the same on forks and non-forks, the average uncertainty of futures rather than where they branch (Fig. 1b). (Per-step noise suits high-dimensional uncertainty, as in weather [11].) The deterministic model's tiny spread is aligned with the branch axis and 32× larger on forks, suggesting, as in the IPE, that uncertainty should enter through the initial conditions. In a pre-registered factorial on 160 held-out states, generative models producing whole clips, with an energy-score loss, 4× more fork data or an action-sensitivity loss catch 0–3% of forks, equally isotropic. Putting the uncertainty into the initial conditions works instead: a learned operator perturbs the start latent by sampled perception error and deterministic dynamics roll it out, catching 52% of forks at 7% false alarms, and 69% at 8% with 4× more fork

Table 2: World models on test set C; spread diagnostics on held-out states with 64 real futures each.

rollouts from	caught	false	spread	main axis	fork ratio
reference (MuJoCo; real futures)	88%	7%	1.00	0.57	4.8
token world model, stochastic	18%	2%	0.98	0.15	1.3
token world model, deterministic	60%	10%	0.04	0.69	32
5 generative variants (Sec. 3.2)	0–3%†	0–1%†	1.1–1.6	0.07–0.10	–
uncertain starts + deterministic, 4× forks	69%†	8%†	0.35	0.48	–

† on 160 held-out training-domain states (75 forks), not set C; none passed the gate to run on a fresh set.

data. Giving a generative model the same perturbed starts does not help (0%): noise in the learned dynamics erases forks. The deterministic recipe plateaus at 68–72% with 3× training, more fork data, frame history or DINO-WM's predictor [12].

4 Related work

Simulation under uncertainty. Beyond the IPE [13], robotics runs Monte Carlo rollouts over pose, friction or model parameters to score grasps and pushes [14, 15] or to build safety filters [16, 17]. Cannizzaro et al. [18] sample measured perception and placement noise through a simulator to predict whether a block tower stands, and Goyal [19] measure offline how fast perturbed chunks diverge. These score a task predicate, labels [20] or a growth rate. Closest in spirit, Zapolsky and Drumwright [21] sample model and sensor parameters to find divergent particle traces when validating robot plans offline; they propose detecting divergence without a goal but, in their experiments, detect it through task failure or by inspection. Our probe instead runs before every chunk, samples the robot's own measured execution and perception errors, and applies an automatic test for whether the outcomes split, a Monte Carlo test of proximity to a basin boundary [22].

Uncertainty-aware monitors. Latent safety filters (also on Jenga) and failure detectors use failure classifiers, epistemic or policy uncertainty [23, 5, 6, 24]: what a model does not know, not whether the physics is at a branch.

World models and their uncertainty. Learned simulators and latent world models [25, 26, 12, 27, 28] struggle with stiff contact [29]; audits find weak physical grounding [30, 31, 32]. Yu et al. [33] propose evaluating world models on interventional action branches, by prediction error and action-effect agreement per branch; our models predict individual futures well yet lose the forks, so a test of branching is needed alongside. Stochastic and ensemble models [34, 35] are evaluated by likelihood, calibration or return, none of which checks whether their spread splits where the dynamics branch.

5 Conclusion and limitations

A robot can judge fragility as people do, by propagating its own measured uncertainty through a rough simulator. For safe deployment: physical uncertainty must be propagated, not just estimated; without a simulator this needs world models with branching fidelity, which our generative models lack despite low average error and uncertain starts partly recover, so benchmarks should measure it; and until learned dynamics have it, our results favour the IPE's design, learned perception and measured uncertainty feeding an approximate physics engine whose scene (shapes, masses, friction) must still be estimated from sensors.

Limitations. Simulation, one task with three blocks; the engine probe uses the simulator's true shapes, masses and friction; one Gaussian models perception error for all images; test sets differ in difficulty; each world model was trained once; pretrained world models (V-JEPA 2-AC, Cosmos) are untested; and the probe flags forks, not certain failures.

Use of generative AI. Generative AI tools (Claude) assisted with code, figures, literature search and drafting. The authors designed the study, checked all results and references, and take full responsibility for the content.

References

[1] P. W. Battaglia, J. B. Hamrick, and J. B. Tenenbaum. Simulation as an engine of physical scene understanding. *Proceedings of the National Academy of Sciences*, 110(45):18327–18332, 2013.

[2] C. Chi, S. Feng, Y. Du, Z. Xu, E. Cousineau, B. Burchfiel, and S. Song. Diffusion policy: Visuomotor policy learning via action diffusion. In *Robotics: Science and Systems*, 2023.

[3] T. Z. Zhao, V. Kumar, S. Levine, and C. Finn. Learning fine-grained bimanual manipulation with low-cost hardware. In *Robotics: Science and Systems*, 2023.

[4] K. Black, N. Brown, D. Driess, A. Esmail, M. Equi, C. Finn, N. Fusai, L. Groom, K. Hausman, B. Ichter, S. Jakubczak, T. Jones, L. Ke, S. Levine, A. Li-Bell, M. Mothukuri, S. Nair, K. Pertsch, L. X. Shi, J. Tanner, Q. Vuong, A. Walling, H. Wang, and U. Zhilinsky. π_0: A vision-language-action flow model for general robot control. *arXiv:2410.24164*, 2024.

[5] J. Seo, K. Nakamura, and A. Bajcsy. Uncertainty-aware latent safety filters for avoiding out-of-distribution failures. In *Conference on Robot Learning*, 2025.

[6] C. Xu, T. K. Nguyen, E. Dixon, C. Rodriguez, P. Miller, R. Lee, P. Shah, R. Ambrus, H. Nishimura, and M. Itkina. Can we detect failures without failure data? Uncertainty-aware runtime failure detection for imitation learning policies. In *Robotics: Science and Systems*, 2025.

[7] K. A. Smith and E. Vul. Sources of uncertainty in intuitive physics. *Topics in Cognitive Science*, 5(1):185–199, 2013.

[8] A. N. Sanborn, V. K. Mansinghka, and T. L. Griffiths. Reconciling intuitive physics and Newtonian mechanics for colliding objects. *Psychological Review*, 120(2):411–437, 2013.

[9] M. Oquab, T. Darcet, T. Moutakanni, H. Vo, M. Szafraniec, V. Khalidov, P. Fernandez, D. Haziza, F. Massa, A. El-Nouby, M. Assran, N. Ballas, W. Galuba, R. Howes, P.-Y. Huang, S.-W. Li, I. Misra, M. Rabbat, V. Sharma, G. Synnaeve, H. Xu, H. Jegou, J. Mairal, P. Labatut, A. Joulin, and P. Bojanowski. DINOv2: Learning robust visual features without supervision. *Transactions on Machine Learning Research*, 2024.

[10] K. A. Ashman, C. M. Bird, and S. E. Zepf. Detecting bimodality in astronomical datasets. *The Astronomical Journal*, 108:2348–2361, 1994.

[11] I. Price, A. Sanchez-Gonzalez, F. Alet, T. R. Andersson, A. El-Kadi, D. Masters, T. Ewalds, J. Stott, S. Mohamed, P. Battaglia, R. Lam, and M. Willson. Probabilistic weather forecasting with machine learning. *Nature*, 637:84–90, 2025.

[12] G. Zhou, H. Pan, Y. LeCun, and L. Pinto. DINO-WM: World models on pre-trained visual features enable zero-shot planning. In *International Conference on Machine Learning*, 2025.

[13] L. Zhou, K. A. Smith, J. B. Tenenbaum, and T. Gerstenberg. Mental Jenga: A counterfactual simulation model of causal judgments about physical support. *Journal of Experimental Psychology: General*, 152(8):2237–2269, 2023.

[14] J. Mahler, J. Liang, S. Niyaz, M. Laskey, R. Doan, X. Liu, J. Aparicio Ojea, and K. Goldberg. Dex-Net 2.0: Deep learning to plan robust grasps with synthetic point clouds and analytic grasp metrics. In *Robotics: Science and Systems*, 2017.

[15] W. C. Agboh and M. R. Dogar. Pushing fast and slow: Task-adaptive planning for non-prehensile manipulation under uncertainty. In *Workshop on the Algorithmic Foundations of Robotics (WAFR)*, pages 160–176, 2018.

[16] A. Johansson, D. Lindmark, V. Wiberg, and M. Servin. Safety filtering of robotic manipulation under environment uncertainty: a computational approach. *arXiv:2509.12674*, 2025.

[17] G. Wang, K. Ren, A. S. Morgan, and K. Hang. ManiDreams: An open-source library for robust object manipulation via uncertainty-aware task-specific intuitive physics. *arXiv:2603.18336*, 2026.

[18] R. Cannizzaro, M. Groom, J. Routley, R. O. Ness, and L. Kunze. COBRA-PPM: A causal Bayesian reasoning architecture using probabilistic programming for robot manipulation under uncertainty. *arXiv:2403.14488*, 2024.

[19] A. Goyal. Measuring the stability assumption behind action chunking. *arXiv:2610.01626*, 2026.

[20] A. Lerer, S. Gross, and R. Fergus. Learning physical intuition of block towers by example. In *International Conference on Machine Learning*, 2016.

[21] S. Zapolsky and E. Drumwright. Particle traces for detecting divergent robot behavior. *arXiv:1608.01606*, 2016.

[22] P. J. Menck, J. Heitzig, N. Marwan, and J. Kurths. How basin stability complements the linear-stability paradigm. *Nature Physics*, 9:89–92, 2013.

[23] K. Nakamura, L. Peters, and A. Bajcsy. Generalizing safety beyond collision-avoidance via latent-space reachability analysis. In *Robotics: Science and Systems*, 2025.

[24] C. Agia, R. Sinha, J. Yang, Z.-a. Cao, R. Antonova, M. Pavone, and J. Bohg. Unpacking failure modes of generative policies: Runtime monitoring of consistency and progress. In *Conference on Robot Learning*, 2024.

[25] A. Sanchez-Gonzalez, J. Godwin, T. Pfaff, R. Ying, J. Leskovec, and P. W. Battaglia. Learning to simulate complex physics with graph networks. In *International Conference on Machine Learning*, 2020.

[26] K. R. Allen, Y. Rubanova, T. Lopez-Guevara, W. Whitney, A. Sanchez-Gonzalez, P. Battaglia, and T. Pfaff. Learning rigid dynamics with face interaction graph networks. In *International Conference on Learning Representations*, 2023.

[27] M. Assran, A. Bardes, D. Fan, Q. Garrido, R. Howes, M. Komeili, M. Muckley, A. Rizvi, C. Roberts, K. Sinha, A. Zholus, S. Arnaud, A. Gejji, A. Martin, F. R. Hogan, D. Dugas, P. Bojanowski, V. Khalidov, P. Labatut, F. Massa, M. Szafraniec, K. Krishnakumar, Y. Li, X. Ma, S. Chandar, F. Meier, Y. LeCun, M. Rabbat, and N. Ballas. V-JEPA 2: Self-supervised video models enable understanding, prediction and planning. *arXiv:2506.09985*, 2025.

[28] D. Hafner, J. Pasukonis, J. Ba, and T. Lillicrap. Mastering diverse control tasks through world models. *Nature*, 640:647–653, 2025.

[29] M. Parmar, M. Halm, and M. Posa. Fundamental challenges in deep learning for stiff contact dynamics. In *IEEE/RSJ International Conference on Intelligent Robots and Systems*, 2021.

[30] D. M. Bear, E. Wang, D. Mrowca, F. J. Binder, H.-Y. F. Tung, R. T. Pramod, C. Holdaway, S. Tao, K. Smith, F.-Y. Sun, L. Fei-Fei, N. Kanwisher, J. B. Tenenbaum, D. L. K. Yamins, and J. E. Fan. Physion: Evaluating physical prediction from vision in humans and machines. In *NeurIPS Datasets and Benchmarks*, 2021.

[31] S. Motamed, L. Culp, K. Swersky, P. Jaini, and R. Geirhos. Do generative video models understand physical principles? In *IEEE/CVF Winter Conference on Applications of Computer Vision*, 2026.

[32] T. Yang, Z. Shen, Z. Mi, Z. Zhang, J. Zhou, J. Ji, J. Dai, J. Chen, B. Chen, and Y. Yang. MiraBench: Evaluating action-conditioned reliability in robotic world models. *arXiv:2605.29360*, 2026.

[33] Y. Yu, S. Zhang, Y. Sheng, H. Ren, and H. Lin. How should world models be evaluated for embodied decision-making? A decision-making-centric position. *arXiv:2606.15032*, 2026.

[34] K. Chua, R. Calandra, R. McAllister, and S. Levine. Deep reinforcement learning in a handful of trials using probabilistic dynamics models. In *Advances in Neural Information Processing Systems*, 2018.

[35] B. Lakshminarayanan, A. Pritzel, and C. Blundell. Simple and scalable predictive uncertainty estimation using deep ensembles. In *Advances in Neural Information Processing Systems*, 2017.