# Branching Fidelity

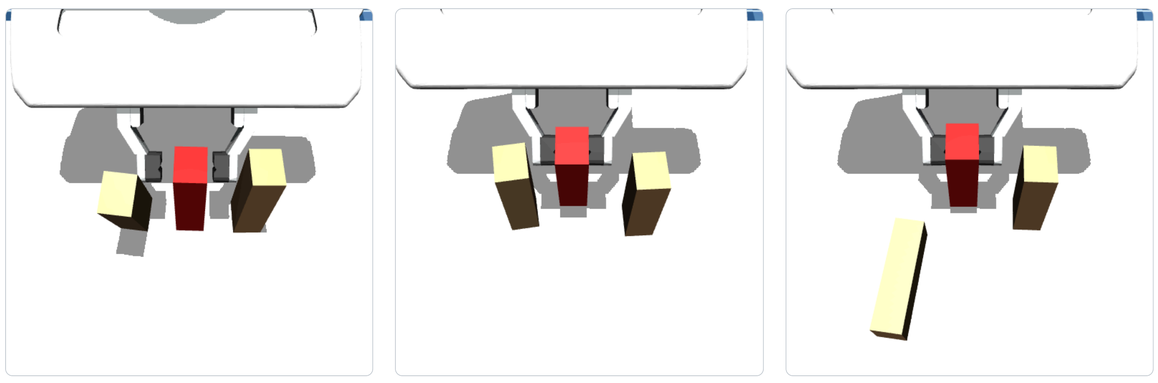
## Testing World Models as Intuitive Physics Engines for Robot Safety

Alikasim Budhwani · Soo Jeon
Mechanical and Mechatronics Engineering, University of Waterloo · a2budhwa@uwaterloo.ca
CoRL 2026 Workshop on Bringing Physics Simulation and World Models Together for Robotic Manipulation · Austin, TX

**Many manipulation tasks hinge on millimetres. A world model useful for them must branch where physics branches, and average prediction error does not test that. The fork probe does.**

## 1 Small errors, different outcomes

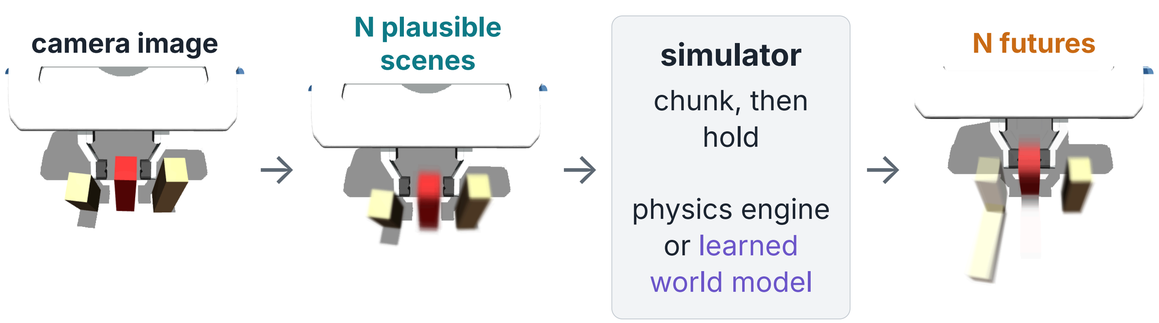
Picking from a stack, pushing a tall bottle, inserting a peg: **a few millimetres of error in the start or the action decide the outcome.** Here one command lifts the middle block cleanly or tips a neighbour.



the chunk starts | half the futures: neighbour stays | the other half: neighbour topples

A state whose futures, sampled over the robot's own small errors, split into different outcomes is a **fork**. A world model's **branching fidelity**: does it fork where physics does?

## 2 The fork probe: a branching-fidelity test



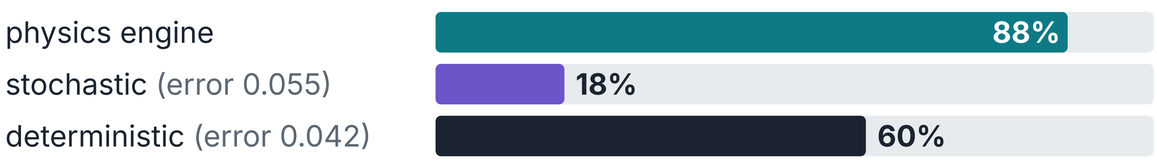

- Samples the robot's **measured** perception (2.5 mm) and execution (~3 mm) errors.
- Flags **two separate groups**, not just spread: on the endings' main axis, two Gaussians must fit better than one (BIC) with a clear gap (Ashman's D > 2), at least 2 futures in the smaller group, at two hold times.
- No task labels, failure data or tuned thresholds.

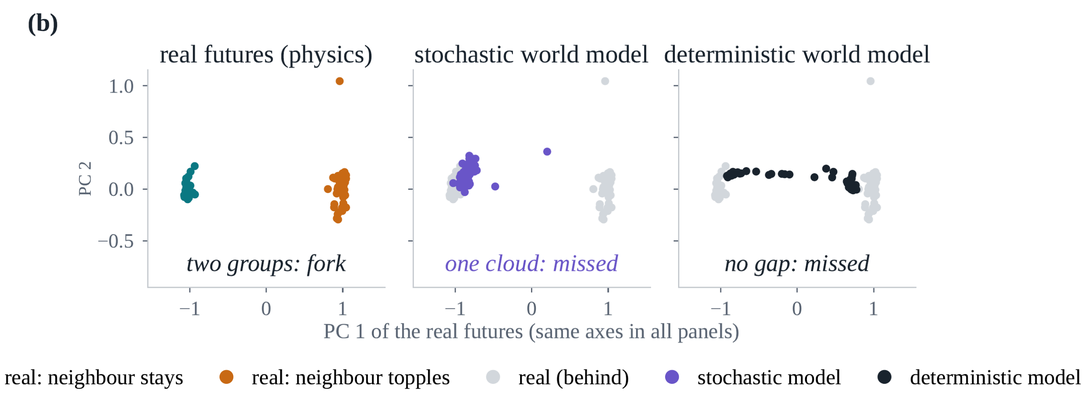
**Upper reference: the physics engine.** MuJoCo with the true masses and friction is not deployable, but sets the bar: **88%** of forks caught, **7%** false alarms, unseen layouts, one camera image.

## 3 Similar error, different branching

Two world models in DINOv2 feature space. By average error they look alike, both **3–4× better than predicting that nothing moves**. The fork probe separates them. Same latent for both (it keeps 97% of real forks), so only where the noise enters differs.

FORKS CAUGHT · SAME UNSEEN TEST SET




## 4 Why

In physics, futures differ only because starts and commands differ, and diverge only near the boundary between outcomes.

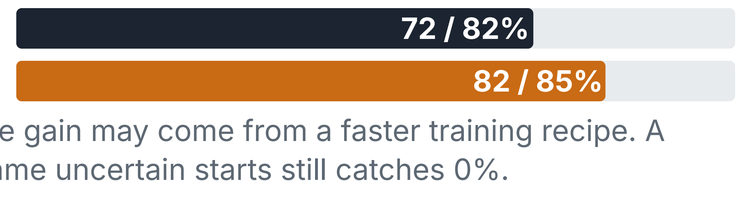
**Stochastic: noise in the wrong place**
Fresh noise at every step: the right *amount* of spread, but in every direction, almost as much on safe states as on forks (1.3× vs 4.8× in physics).

**Deterministic: right place, no gap**
Spread comes only from perturbed inputs, as in physics: along the branch, 32× larger on forks. But squared-error training predicts the average, blurring the branches.

## 5 Fork-aware training

**Uncertain starts:** put the uncertainty where it is, in the robot's knowledge of the scene; perturb the start, roll out deterministically. **Curate with the split test:** find each training scene's rare branch without labels and oversample it.

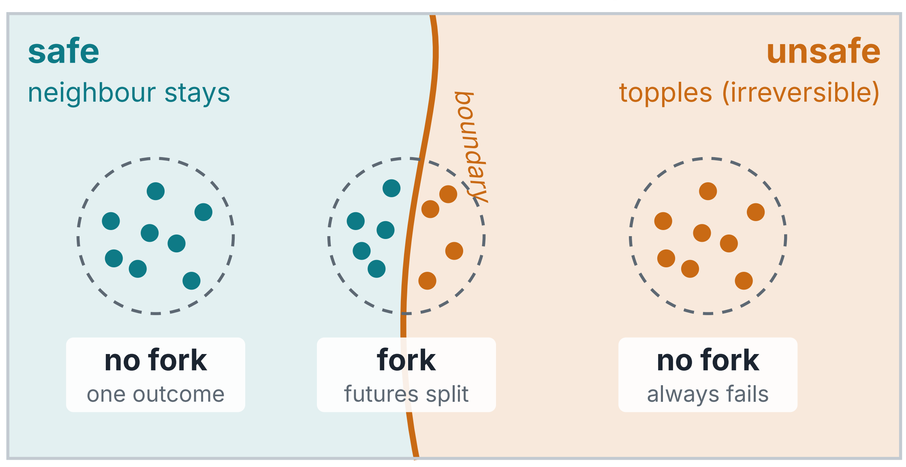
FORKS CAUGHT · TWO FRESH TEST SETS




False alarms 7–11%. Part of the gain may come from a faster training recipe. A generative model given the same uncertain starts still catches 0%.

## 6 Application: a label-free safety signal

People judge whether a stack will fall by **imagining it a few times** with slightly different guesses of where things are; if the outcomes disagree, it is precarious (the **intuitive physics engine** of cognitive science). The probe is this check with measured errors. When one branch is irreversible, as for toppling rigid bodies, a fork marks the **edge of the safe set**.




A proxy, not a failure detector: two harmless branches give a false alarm; an action that always fails, or gradual harm (spills), never splits. Our follow-up builds a label-free safety filter on it.

## 7 Takeaways

**Report branching fidelity**
beside average error: does the model fork where physics does?

**Uncertainty in the start**
Deterministic dynamics from perturbed starts keep the branch; per-step noise smears it.

**Train for rare branches**
The probe's own split test finds them without labels.

**Forks flag the safe-set edge**
in rigid-body tasks; pair them with a harm signal.

Limitations: simulation, one three-block task; the probe and its ground truth assume the same measured error model, so a real-robot check is needed; the split test sees at most two groups along one axis; small world models, trained once.